



Using Pin for Compiler and Computer Architecture Research and Education

PLDI Tutorial
Sunday, June 10, 2007

Kim Hazelwood
David Kaeli
Dan Connors
Vijay Janapa Reddi



Part One: Introduction and Overview

Kim Hazelwood
David Kaeli
Dan Connors
Vijay Janapa Reddi

What is Instrumentation?



A technique that inserts extra code into a program to collect runtime information

Instrumentation approaches:

- Source instrumentation:
 - Instrument source programs
- **Binary instrumentation:**
 - Instrument executables directly

Why use Dynamic Instrumentation?



- ✓ No need to recompile or relink
- ✓ Discover code at runtime
- ✓ Handle dynamically-generated code
- ✓ Attach to running processes

How is Instrumentation used in Compiler Research?



Program analysis

- Code coverage
- Call-graph generation
- Memory-leak detection
- Instruction profiling

Thread analysis

- Thread profiling
- Race detection

How is Instrumentation used in Computer Architecture Research?



- Trace Generation
- Branch Predictor and Cache Modeling
- Fault Tolerance Studies
- Emulating Speculation
- Emulating New Instructions

Advantages of Pin Instrumentation



Easy-to-use Instrumentation:

- Uses dynamic instrumentation
 - Do not need source code, recompilation, post-linking

Programmable Instrumentation:

- Provides rich APIs to write in C/C++ your own instrumentation tools (called Pintools)

Multiplatform:

- Supports x86, x86-64, Itanium, Xscale
- Supports Linux, Windows, MacOS

Robust:

- Instruments real-life applications: Database, web browsers, ...
- Instruments multithreaded applications
- Supports signals

Efficient:

- Applies compiler optimizations on instrumentation code

Other Advantages



- **Robust and stable**
 - Pin can run itself!
 - 12+ active developers
 - Nightly testing of 25000 binaries on 15 platforms
 - Large user base in academia and industry
 - Active mailing list (Pinheads)
- **14,000 downloads**

Using Pin



Launch and instrument an application

```
$ pin -t pintool -- application
```

↑
Instrumentation engine
(provided in the kit)

↑
Instrumentation tool
(write your own, or use one
provided in the kit)

Attach to and instrument an application

```
$ pin -t pintool -pid 1234
```

Pin Instrumentation APIs



Basic APIs are architecture independent:

- Provide common functionalities like determining:
 - Control-flow changes
 - Memory accesses

Architecture-specific APIs

- e.g., Info about segmentation registers on IA32

Call-based APIs:

- Instrumentation routines
- Analysis routines

Instrumentation vs. Analysis



Concepts borrowed from the ATOM tool:

Instrumentation routines define where instrumentation is **inserted**

- e.g., before instruction
- ☞ Occurs *first time* an instruction is executed

Analysis routines define what to do when instrumentation is **activated**

- e.g., increment counter
- ☞ Occurs *every time* an instruction is executed

Pintool 1: Instruction Count



```
sub $0xff, %edx
counter++;
cmp %esi, %edx
counter++;
jle <L1>
counter++;
mov $0x1, %edi
counter++;
add $0x10, %eax
counter++;
```

Pintool 1: Instruction Count Output



```
$ /bin/ls
```

```
Makefile imageload.out itrace proccount
imageload inscount0 atrace itrace.out
```

```
$ pin -t inscount0 -- /bin/ls
```

```
Makefile imageload.out itrace proccount
imageload inscount0 atrace itrace.out
```

```
Count 422838
```


ManualExamples/inscount0.cpp



```
#include <iostream>
#include "pin.h"

UINT64 icount = 0;

void docount() { icount++; }
void Instruction(INS ins, void *v)
{
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)docount, IARG_END);
}

void Fini(INT32 code, void *v)
{
    std::cerr << "Count " << icount << endl;
}

int main(int argc, char * argv[])
{
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

analysis routine

instrumentation routine

Pintool 2: Instruction Trace



```
Print(ip);
sub $0xff, %edx
Print(ip);
cmp %esi, %edx
Print(ip);
jle <L1>
Print(ip);
mov $0x1, %edi
Print(ip);
add $0x10, %eax
```

Need to pass ip argument to the analysis routine (printip())

Pintool 2: Instruction Trace Output



```
$ pin -t itrace -- /bin/ls
Makefile imageload.out itrace proccount
imageload inscount0 atrace itrace.out
```

```
$ head -4 itrace.out
0x40001e90
0x40001e91
0x40001ee4
0x40001ee5
```

ManualExamples/itrace.cpp



```
#include <stdio.h>
#include "pin.H"
FILE * trace;

void printip(void *ip) { fprintf(trace, "%p\n", ip); }
                                analysis routine

void Instruction(INS ins, void *v) {
                                instrumentation routine
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)printip,
                  IARG_INST_PTR, IARG_END);
}

void Fini(INT32 code, void *v) { fclose(trace); }

int main(int argc, char * argv[]) {
    trace = fopen("itrace.out", "w");
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);

    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

Examples of Arguments to Analysis Routine



IARG_INST_PTR

- Instruction pointer (program counter) value

IARG_UINT32 <value>

- An integer value

IARG_REG_VALUE <register name>

- Value of the register specified

IARG_BRANCH_TARGET_ADDR

- Target address of the branch instrumented

IARG_MEMORY_READ_EA

- Effective address of a memory read

And many more ... (refer to the Pin manual for details)

Instrumentation Points



Instrument points relative to an instruction:

- Before (IPOINT_BEFORE)
- After:
 - Fall-through edge (IPOINT_AFTER)
 - Taken edge (IPOINT_TAKEN_BRANCH)

```
count() → cmp    %esi, %edx    count()
           jle    <L1>         ↘
count() → mov    $0x1, %edi     <L1>: mov    $0x8,%edi
```

Instrumentation Granularity



Instrumentation can be done at three different granularities:

- Instruction
- Basic block
 - A sequence of instructions terminated at a control-flow changing instruction
 - Single entry, single exit
- Trace
 - A sequence of basic blocks terminated at an unconditional control-flow changing instruction
 - Single entry, multiple exits

```
sub    $0xff, %edx  
cmp    %esi, %edx  
jle    <L1>
```

```
mov    $0x1, %edi  
add    $0x10, %eax  
jmp    <L2>
```

1 Trace, 2 BBs, 6 insts

Recap of Pintool 1: Instruction Count



```
counter++;  
sub    $0xff, %edx  
counter++;  
cmp    %esi, %edx  
counter++;  
jle    <L1>  
counter++;  
mov    $0x1, %edi  
counter++;  
add    $0x10, %eax
```

Straightforward, but the counting can be more efficient

Pintool 3: Faster Instruction Count



```
counter += 3
sub  $0xff, %edx

cmp  %esi, %edx

jle  <L1>

counter += 2
mov  $0x1, %edi

add  $0x10, %eax
```

basic blocks (bbl)

```
#include <stdio.h>
#include "pin.H"
UINT64 icount = 0;
void docount(INT32 c) { icount += c; }
void Trace(TRACE trace, void *v) {
    for (BBL bbl = TRACE_BblHead(trace);
         BBL_Valid(bbl); bbl = BBL_Next(bbl)) {
        BBL_InsertCall(bbl, IPOINT_BEFORE, (AFUNPTR)docount,
                       IARG_UINT32, BBL_NumIns(bbl), IARG_END);
    }
}
void Fini(INT32 code, void *v) {
    fprintf(stderr, "Count %lld\n", icount);
}
int main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    TRACE_AddInstrumentFunction(Trace, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

ManualExamples/inscount1.cpp



analysis routine

instrumentation routine

Modifying Program Behavior



Pin allows you not only to observe but also change program behavior

Ways to change program behavior:

- Add/delete instructions
- Change register values
- Change memory values
- Change control flow

Instrumentation Library



```
#include <iostream>
#include "pin.H"

UINT64 icount = 0;

VOID Fini(INT32 code, VOID *v) {
    std::cerr << "Count " << icount << endl;
}

VOID dcount() {
    icount++;
}

VOID Instruction(INS ins, VOID *v) {
    INS_InsertCall(ins, IPOINT_BEFORE,
    (IOPICALLBACK) dcount, 0, 0);
}

int main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

Instruction counting Pin Tool

```
#include <iostream>
#include "pin.H"
#include "instlib.H"

INSTLIB::ICOUNT icount;

VOID Fini(INT32 code, VOID *v) {
    cout << "Count" << icount.Count() << endl;
}

int main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    PIN_AddFiniFunction(Fini, 0);
    icount.Activate();
    PIN_StartProgram();
    return 0;
}
```

Useful InstLib abstractions



- **ICOUNT**
 - # of instructions executed
- **FILTER**
 - Instrument specific routines or libraries only
- **ALARM**
 - Execution count timer for address, routines, etc.
- **FOLLOW_CHILD**
 - Inject Pin into new process created by parent process
- **TIME_WARP**
 - Preserves RDTSC behavior across executions
- **CONTROL**
 - Limit instrumentation address ranges

Debugging Pintools



1. Invoke gdb with your pintool (don't "run")

```
$ gdb inscount0  
(gdb)
```

2. In another window, start your pintool with the "-pause_tool" flag

```
$ pin -pause_tool 5 -t inscount0 -- /bin/ls  
Pausing to attach to pid 32017
```

3. Go back to gdb window:

- a) Attach to the process
- b) "cont" to continue execution; can set breakpoints as usual

```
(gdb) attach 32017  
(gdb) break main  
(gdb) cont
```



Pin Internals

Pin Source Code Organization

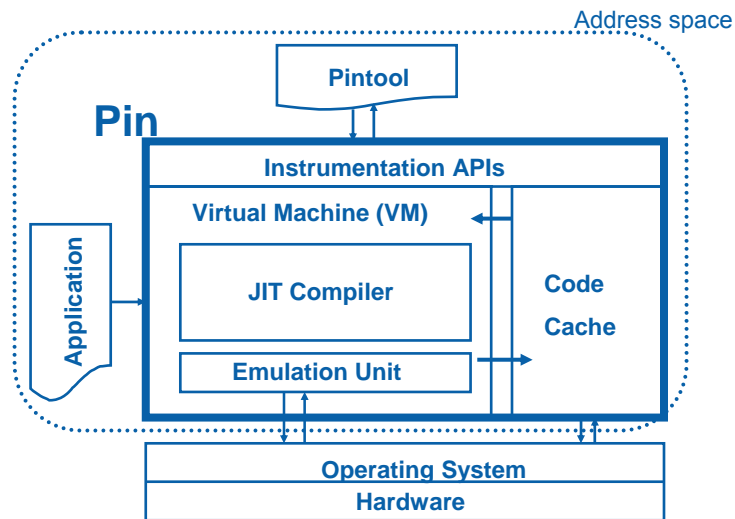


Pin source organized into generic, architecture-dependent, OS-dependent modules:

Architecture	#source files	#source lines
Generic	87 (48%)	53595 (47%)
x86 (32-bit+ 64-bit)	34 (19%)	22794 (20%)
Itanium	34 (19%)	20474 (18%)
ARM	27 (14%)	17933 (15%)
TOTAL	182 (100%)	114796 (100%)

 ~50% code shared among architectures

Pin's Software Architecture




31

Pin PLDI Tutorial 2007


Dynamic Instrumentation

Original code



Pin fetches trace starting block 1 and start instrumentation

Code cache




32

Pin PLDI Tutorial 2007


Dynamic Instrumentation



Original code



Code cache



Pin transfers control into code cache (block 1)


33

Pin PLDI Tutorial 2007


Dynamic Instrumentation



Original code



Code cache



Pin fetches and instrument a new trace

34

Pin PLDI Tutorial 2007

Implementation Challenges



- **Linking**
 - Straightforward for direct branches
 - Tricky for indirects, invalidations
- **Re-allocating registers**
- **Maintaining transparency**
- **Self-modifying code**
- **Supporting MT applications...**

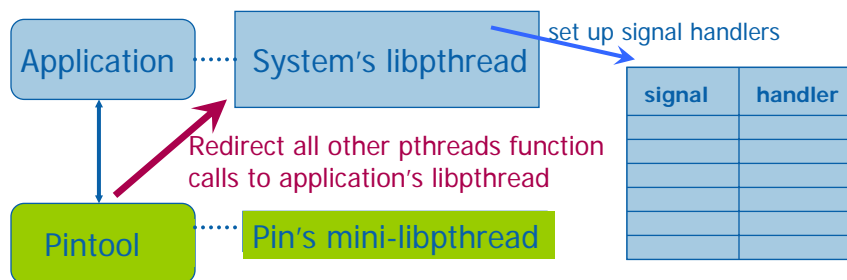
Pin's Multithreading Support




Thread-safe accesses Pin, Pintool, and App

- Pin: One thread in the VM at a time
- Pintool: Locks, ThreadID, event notification
- App: Thread-local spill area

Providing pthreads functions to instrumentation tools




Pin Overhead



37

Pin PLDI Tutorial 2007

Adding User Instrumentation



38

Pin PLDI Tutorial 2007



Optimizing Pintools

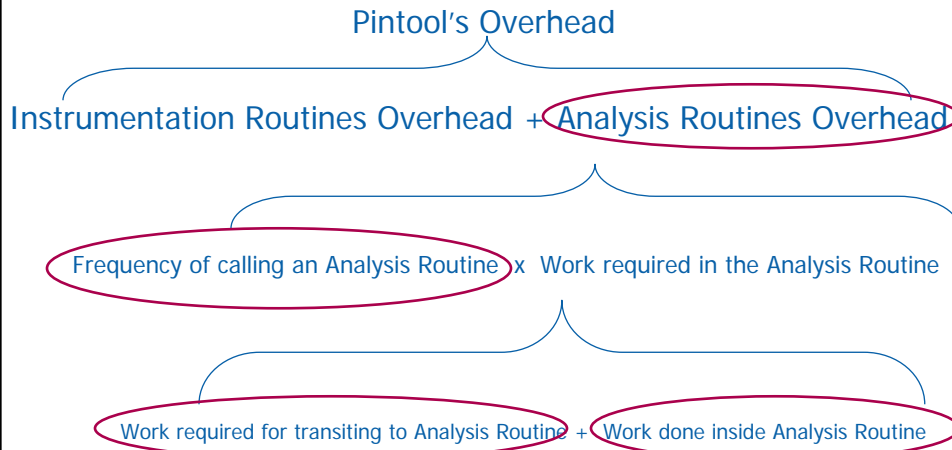
Reducing Instrumentation Overhead

Total Overhead = Pin Overhead + Pintool Overhead

- Pin team's job is to minimize this
- ~5% for SPECfp and ~20% for SPECint

- Pintool writers can help minimize this!

Reducing the Pintool's Overhead



Analysis Routines: Reduce Call Frequency



Key: Instrument at the largest granularity whenever possible

Trace > Basic Block > Instruction

Slower Instruction Counting



```
counter++;  
sub $0xff, %edx  
counter++;  
cmp %esi, %edx  
counter++;  
jle <L1>  
counter++;  
mov $0x1, %edi  
counter++;  
add $0x10, %eax
```

Faster Instruction Counting



Counting at BBL level

```
counter += 3  
sub $0xff, %edx  
  
cmp %esi, %edx  
  
jle <L1>
```

```
counter += 2  
mov $0x1, %edi  
  
add $0x10, %eax
```

Counting at Trace level

```
counter += 5  
sub $0xff, %edx  
  
cmp %esi, %edx  
  
jle <L1>
```

```
mov $0x1, %edi  
  
add $0x10, %eax
```

counter -= 2

L1

Reducing Work in Analysis Routines



Key: Shift computation from analysis routines to instrumentation routines whenever possible

Edge Counting: a Slower Version



```
...
void docount2(ADDRINT src, ADDRINT dst, INT32 taken)
{
    COUNTER *pedg = Lookup(src, dst);
    pedg->count += taken;
}
void Instruction(INS ins, void *v) {
    if (INS_IsBranchOrCall(ins))
    {
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)docount2,
                       IARG_INST_PTR, IARG_BRANCH_TARGET_ADDR,
                       IARG_BRANCH_TAKEN, IARG_END);
    }
}
...
```


Edge Counting: a Faster Version



```
void docount(COUNTER* pedg, INT32 taken) {
    pedg->count += taken;
}

void docount2(ADDRINT src, ADDRINT dst, INT32 taken) {
    COUNTER *pedg = Lookup(src, dst);
    pedg->count += taken;
}

void Instruction(INS ins, void *v) {
    if (INS_IsDirectBranchOrCall(ins)) {
        COUNTER *pedg = Lookup(INS_Address(ins),
                               INS_DirectBranchOrCallTargetAddress(ins));
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) docount,
                       IARG_ADDRINT, pedg, IARG_BRANCH_TAKEN, IARG_END);
    } else
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) docount2,
                       IARG_INST_PTR, IARG_BRANCH_TARGET_ADDR,
                       IARG_BRANCH_TAKEN, IARG_END);
}
...
```

Reducing Work for Analysis Transitions



Key: Help Pin's optimizations apply to your analysis routines:

- Inlining
- Scheduling

Inlining



Inlinable

```
int docount0(int i) {  
    x[i]++;  
    return x[i];  
}
```

Not-inlinable

```
int docount1(int i) {  
    if (i == 1000)  
        x[i]++;  
    return x[i];  
}
```

Not-inlinable

```
int docount2(int i) {  
    x[i]++;  
    printf("%d", i);  
    return x[i];  
}
```

Not-inlinable

```
void docount3() {  
    for(i=0;i<100;i++)  
        x[i]++;  
}
```

Conditional Inlining



Inline a common scenario where the analysis routine has a single "if-then"

- The "If" part is always executed
- The "then" part is rarely executed

Pintool writer breaks such an analysis routine into two:

- **INS_InsertIfCall** (ins, ..., (AFUNPTR)doif, ...)
- **INS_InsertThenCall** (ins, ..., (AFUNPTR)dothen, ...)

IP-Sampling (a Slower Version)



```
const INT32 N = 10000; const INT32 M = 5000;
INT32 icount = N;

VOID IpSample(VOID* ip) {
    --icount;
    if (icount == 0) {
        fprintf(trace, "%p\n", ip);
        icount = N + rand()%M; //icount is between <N, N+M>
    }
}

VOID Instruction(INS ins, VOID *v) {
    INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)IpSample,
        IARG_INST_PTR, IARG_END);
}
```

IP-Sampling (a Faster Version)



```
INT32 CountDown() {
    --icount;
    return (icount==0);
}
VOID PrintIp(VOID *ip) {
    fprintf(trace, "%p\n", ip);
    icount = N + rand()%M; //icount is between <N, N+M>
}

VOID Instruction(INS ins, VOID *v) {
    // CountDown() is always called before an inst is executed
    INS_InsertIfCall(ins, IPOINT_BEFORE, (AFUNPTR)CountDown,
        IARG_END);

    // PrintIp() is called only if the last call to CountDown()
    // returns a non-zero value
    INS_InsertThenCall(ins, IPOINT_BEFORE, (AFUNPTR)PrintIp,
        IARG_INST_PTR, IARG_END);
}
```

} inlined

} not inlined

Instrumentation Scheduling



If an instrumentation can be inserted anywhere in a basic block:

- Let Pin know via **IPOINT_ANYWHERE**
- Pin will find the best point to insert the instrumentation to minimize register spilling

```
#include <stdio.h>
#include "pin.H"
UINT64 icount = 0;
void docount(INT32 c) { icount += c; }
void Trace(TRACE trace, void *v) {
    for (BBL bbl = TRACE_BblHead(trace);
         BBL_Valid(bbl); bbl = BBL_Next(bbl)) {
        BBL_InsertCall(bbl, IPOINT_ANYWHERE, (AFUNPTR)docount,
                       IARG_UINT32, BBL_NumIns(bbl), IARG_END);
    }
}
void Fini(INT32 code, void *v) {
    fprintf(stderr, "Count %lld\n", icount);
}
int main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    TRACE_AddInstrumentFunction(Trace, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

ManualExamples/inscount1.cpp



analysis routine

instrumentation routine

Conclusions



A dynamic instrumentation system for building your own program analysis tools

Runs on multiple platforms:

- IA-32, Intel64, Itanium, and XScale
- Linux, Windows, MacOS

Works on real-life applications

Efficient instrumentation (especially with your help!)

Part Two: Fundamental Concepts in Compilers and Architecture

Kim Hazelwood

David Kaeli

Dan Connors

Vijay Janapa Reddi

Pin Applications



Sample tools in the Pin distribution:

- Cache simulators, branch predictors, address tracer, syscall tracer, edge profiler, stride profiler

Some tools developed and used inside Intel:

- *Opcodemix* (analyze code generated by compilers)
- *PinPoints* (find representative regions in programs to simulate)
- A tool for detecting memory bugs

Companies are writing their own Pintools

Universities use Pin in teaching and research

Tools for Program Analysis



Debugtrace – debugging/program understanding aid, can see general call traces, instruction traces, includes reads and writes of registers and memory

Malloctrace – traces of execution of specific functions

Insmix – statistics/characterization

Statica – static analysis of binaries

Compiler Bug Detection



Opcodemix uncovered a compiler bug for crafty

Instruction Type	Compiler A Count	Compiler B Count	Delta	
*total	712M	618M	-94M	←
XORL	94M	94M	0M	
TESTQ	94M	94M	0M	
RET	94M	94M	0M	
PUSHQ	94M	0M	-94M	←
POPQ	94M	0M	-94M	←
JE	94M	0M	-94M	←
LEAQ	37M	37M	0M	
JNZ	37M	131M	94M	←

Thread Checker Basics



Detect common parallel programming bugs:

- Data races, deadlocks, thread stalls, threading API usage violations

Instrumentation used:

- Memory operations
- Synchronization operations (via function replacement)
- Call stack


Pin-based prototype

- Runs on Linux, x86 and x86_64
- A Pintool ~2500 C++ lines

Thread Checker Results



Potential errors in SPECOMP01 reported by Thread Checker
(4 threads were used)



61

Pin PLDI Tutorial 2007

PinContrib\ThreadChecker\Tests\SPECOMP01-art_m-train2-hackedVia32\ThreadChecker.thr: Diagnostics

Context [Best]	Sev...	Description	Counts	1st Access...	1st Access [Best]	2nd Access [Routine]	2nd Access [Best]
"scanner.c": 1015	12	Memory write of unknown at "scanner.c": 152 conflicts with a prior memory read unknown at "scanner.c": 152 (anti dependence)	44	match	"scanner.c": 152	match	"scanner.c": 152
"scanner.c": 1015	13	Memory read of unknown at "scanner.c": 154 conflicts with a prior memory write of unknown at "scanner.c": 155 (flow dependence)	44	match	"scanner.c": 155	match	"scanner.c": 154
"scanner.c": 1015	14	Memory write of unknown at "scanner.c": 155 conflicts with a prior memory write of unknown at "scanner.c": 155 (output dependence)	44	match	"scanner.c": 155	match	"scanner.c": 155
Group 2: "scanner.c": 1627 (1 item)							
"scanner.c": 1627	0	Function call pthread_mutex_unlock fails with 1 returned at "scanner.c": 1627	2	main	"scanner.c": 1627	main	"scanner.c": 1627

PinContrib\ThreadChecker\Tests\SPECOMP01-art_m-train2-hackedVia32\ThreadChecker.thr (ID=13) 2nd Access

Stack Trace: match "scanner.c": 154

Address	Line	Source
	134	static int num_list_items_added = 0;
	135	
	136	
	137	void add_list_item(int neuron, double vigi, int x, int y)
	138	{
	139	struct lnkNode* ptr;
	140	struct lnkNode* temp;
0x4920	141	temp = (struct lnkNode*) malloc(sizeof(struct lnkNode));
0x492D	142	if (temp == NULL) {
0x4AF7	143	printf("malloc problems\n");
0x4B01	144	exit(1);
	145	}
	146	
0x4935	147	temp->x = x;
0x4939	148	temp->y = y;
0x4957	149	temp->F2Neuron = neuron;
0x493D	150	temp->Vigilance = vigi;
	151	
0x494D	152	++num_list_items_added;
	153	
0x4959	154	temp->next = head;
0x495F	155	head = temp;
	156	

a documented data race in the art benchmark is detected

Instrumentation Driven Simulation



Fast exploratory studies

- Instrumentation \approx native execution
- Simulation speeds at MIPS

Characterize complex applications

- E.g. Oracle, Java, parallel data-mining apps

Simple to build instrumentation tools

- Tools can feed simulation models in real time
- Tools can gather instruction traces for later use

Performance Models



Branch Predictor Models:

- PC of conditional instructions
- Direction Predictor: Taken/not-taken information
- Target Predictor: PC of target instruction if taken


Cache Models:

- Thread ID (if multi-threaded workload)
- Memory address
- Size of memory operation
- Type of memory operation (Read/Write)

Simple Timing Models:

- Latency information

Branch Predictor Model



BPSim Pin Tool

- Instruments all branches
- Uses API to set up call backs to analysis routines

Branch Predictor Model:

- Detailed branch predictor simulator

BP Implementation



```
BranchPredictor myBPU;
```

ANALYSIS

```
VOID ProcessBranch(ADDRINT PC, ADDRINT targetPC, bool BrTaken) {
    BP_Info pred = myBPU.GetPrediction( PC );
    if( pred.Taken != BrTaken ) {
        // Direction Mispredicted
    }
    if( pred.predTarget != targetPC ) {
        // Target Mispredicted
    }
    myBPU.Update( PC, BrTaken, targetPC);
}
```

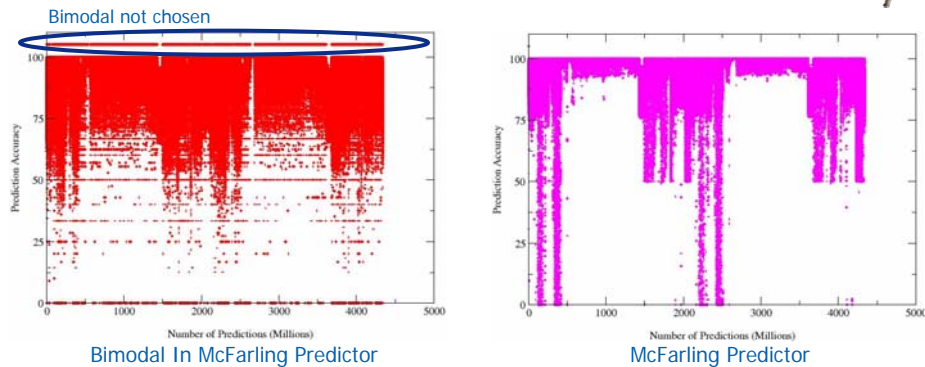
INSTRUMENT

```
VOID Instruction(INS ins, VOID *v)
{
    if( INS_IsDirectBranchOrCall(ins) || INS_HasFallThrough(ins) )
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) ProcessBranch,
            ADDRINT, INS_Address(ins),
            IARG_UINT32, INS_DirectBranchOrCallTargetAddress(ins),
            IARG_BRANCH_TAKEN, IARG_END);
}
```

MAIN

```
int main() {
    PIN_Init();
    INS_AddInstrumentationFunction(Instruction, 0);
    PIN_StartProgram();
}
```

Branch Predictor Performance - GCC



Branch prediction accuracies range from 0-100%

Branches are hard to predict in some phases

- Can simulate these regions alone by fast forwarding to them in real time

Performance Model Inputs



Branch Predictor Models:

- PC of conditional instructions
- Direction Predictor: Taken/not-taken information
- Target Predictor: PC of target instruction if taken


Cache Models:

- Thread ID (if multi-threaded workload)
- Memory address
- Size of memory operation
- Type of memory operation (Read/Write)

Simple Timing Models:

- Latency information

Cache Simulators



Cache Pin Tool

- Instruments all instructions that reference memory
- Use API to set up call backs to analysis routines

Cache Model:

- Detailed cache simulator

Cache Implementation



```

CACHE_t CacheHierarchy[MAX_NUM_THREADS][MAX_NUM_LEVELS];

ANALYSIS
VOID MemRef(int tid, ADDRINT addrStart, int size, int type) {
    for(addr=addrStart; addr<(addrStart+size); addr+=LINE_SIZE)
        LookupHierarchy( tid, FIRST_LEVEL_CACHE, addr, type);
}
VOID LookupHierarchy(int tid, int level, ADDRINT addr, int accessType){
    result = cacheHier[tid][cacheLevel]->Lookup(addr, accessType );
    if( result == CACHE_MISS ) {
        if( level == LAST_LEVEL_CACHE ) return;
        LookupHierarchy(tid, level+1, addr, accessType);
    }
}

INSTRUMENT
VOID Instruction(INS ins, VOID *v)
{
    if( INS_IsMemoryRead(ins) )
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) MemRef,
            IARG_THREAD_ID, IARG_MEMORYREAD_EA, IARG_MEMORYREAD_SIZE,
            IARG_UINT32, ACCESS_TYPE_LOAD, IARG_END);
    if( INS_IsMemoryWrite(ins) )
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) MemRef,
            IARG_THREAD_ID, IARG_MEMORYWRITE_EA, IARG_MEMORYWRITE_SIZE,
            IARG_UINT32, ACCESS_TYPE_STORE, IARG_END);
}

MAIN
int main() {
    PIN_Init();
    INS_AddInstrumentationFunction(Instruction, 0);
    PIN_StartProgram();
}
  
```

Performance Models



Branch Predictor Models:

- PC of conditional instructions
- Direction Predictor: Taken/not-taken information
- Target Predictor: PC of target instruction if taken

Cache Models:

- Thread ID (if multi-threaded workload)
- Memory address
- Size of memory operation
- Type of memory operation (Read/Write)

Simple Timing Models:

- Latency information

Simple Timing Model



Assume 1-stage pipeline

- T_i cycles for instruction execution

Assume branch misprediction penalty

- T_b cycles penalty for branch misprediction


Assume cache access & miss penalty

- T_l cycles for demand reference to cache level l
- T_m cycles for demand reference to memory

$$\text{Total cycles} = \alpha T_i + \beta T_b + \sum_{l=1}^{\text{LLC}} A_l T_l + \eta T_m$$

α = instruction count; β = # branch mispredicts ;
 A_l = # accesses to cache level l ; η = # last level cache (LLC) misses

Performance - GCC



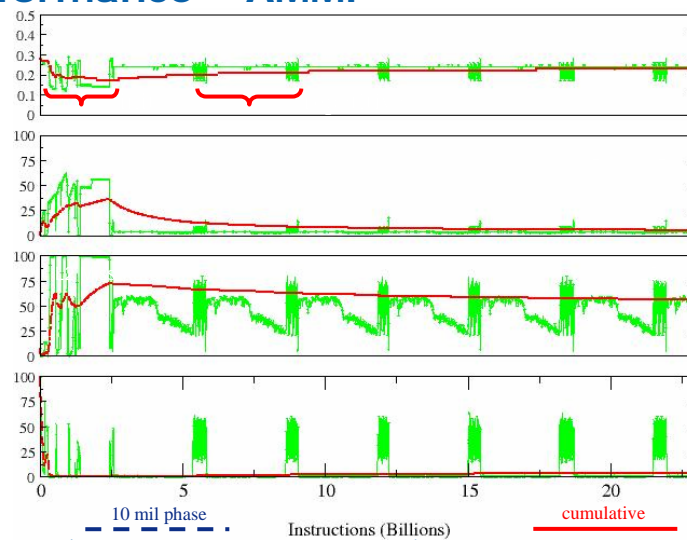
Several phases of execution

- Important to pick the correct phase of execution

73

Pin PLDI Tutorial 2007

Performance – AMMP



One loop (3 billion instructions) is representative

- High miss rate at beginning; exploits locality at end

74

Pin PLDI Tutorial 2007

More Fundamental Concepts using Pin

Kim Hazelwood

David Kaeli

Dan Connors

Vijay Janapa Reddi

Moving from 32-bit to 64-bit Applications


- Intuition would tell us that code expansion will occur
- How can compiler writers exploit the features of a 64-bit ISA?
- What type of programs will benefit from this migration?
- What data types make use of the move to 64 bits?
- How do we begin to identify the reasons for the performance results shown in this table?

➤ Profiling with Pin!


Benchmark	Language	64-bit vs. 32-bit speedup
perlbench	C	3.42%
bzip2	C	15.77%
gcc	C	-18.09%
mcf	C	-26.35%
gobmk	C	4.97%
hmmer	C	34.34%
sjeng	C	14.21%
libquantum	C	35.38%
h264ref	C	35.35%
omnetpp	C++	-7.83%
astar	C++	8.46%
xalancbmk	C++	-13.65%
Average		7.16%

Ye06, IISWC2006

Code Size *Increases* in 64-bit Mode



Dynamic Instruction Count *Decreases* in 64-bit Mode



Observations




Code density increases in 64-bit mode

- More registers help
- 64-bit integer arithmetic helps a lot in the case of libquantum

Concern over the reduction of decoding efficiency in 64-bit is not substantiated

Instruction Cache (L1) Request Rate *Increases in 64-bit Mode*




Instruction Cache Miss Rate Comparison



Code size increases due to:

- Increased instruction length – 10% on average 64-bit mode
- Doubling the size of long and pointer data types




81

Pin PLDI Tutorial 2007

Data Cache (L1) Request Rate



Decreases in 64-bit mode




82

Pin PLDI Tutorial 2007

Data Cache Miss Rate Comparison



Data cache miss rate increases in 64-bit



Observations



The instruction cache miss rate is very low in both 64-bit and 32-bit modes

The data cache request rate decreases significantly in 64-bit mode

- Extra registers help

The data cache miss rate increases in 64-bit mode

- The increased size of long and pointer data types has an adverse impact on data cache performance

Moving from 32-bit to 64-bit Applications



- Common assumptions associated with changes in architecture word sizes need to be studied carefully
- All of these analyses were done with slightly modified versions of the Pintools in the SimpleExamples directory shipped with Pin
 - icount.cpp
 - ilenmix.cpp
 - opcodemix.cpp
 - icache.cpp
 - dcache.cpp

Conclusions



Instrumentation based simulation:

- Simple compared to detailed models
- Can easily run complex applications
- Provides insight on workload behavior over their entire runs in a reasonable amount of time

Illustrated the use of Pin for:

- Compilers
 - Bug detection, thread analysis
- Computer architecture
 - Branch predictors, cache simulators, timing models, architecture width
- Architecture changes
 - Move from 32-bit to 64-bit



Part Three: Advanced Concepts in Compilers and Architecture

Kim Hazelwood

David Kaeli

Dan Connors

Vijay Janapa Reddi

Using Pin in Security Research



- How do we design architectural extensions to accelerate SPAM filtering and anti-virus scanning workloads?
- How do we track dynamic information flow to detect zero-day attack intrusions?

➤ Pin!!

The cost of SPAM – (2006 study)



- Internet users receive 12.4B SPAM email messages daily
- Greater than 40% of all email messages received daily are SPAM (22% for corporate emails)
- SPAM volume is estimated to increase by 63% in 2007
- SPAM filters are the current state-of-the-art in reducing this impact

SPAM Filters



- Server-side filtering typically uses Bayesian classification
- Probability that a document contains SPAM is computed as:

$$Prob(\text{spam} \mid \text{words}) = Prob(\text{word} \mid \text{spam}) * Prob(\text{spam}) / Prob(\text{word})$$

- It is important to train the classifier prior to filtering

$$P(C = \text{spam} \mid X = x) / P(C = \text{ham} \mid X = x) > \lambda$$

where λ represents a threshold (typical: $\lambda=9$ to 999)

What's Hot in Bayesian Classification (250 training messages)



Bogofilter

Function	Calls	Instructions	Total
word_cmp	63,452	155	9,835,060
yylex	19,634	5,209	102,278,715

Spamprobe

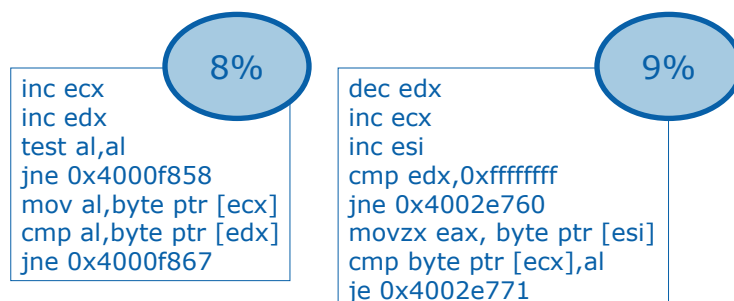
Function	Calls	Instructions	Total
strcmp	1,312,964	11	14,442,604
readLine	152	599	91,048

***String comparison dominates execution**

Bayesian Classification - Bogofilter



- 2 hot code traces dominate (2 basic blocks each) the dynamic execution stream



Bayesian Classification - Spamprobe



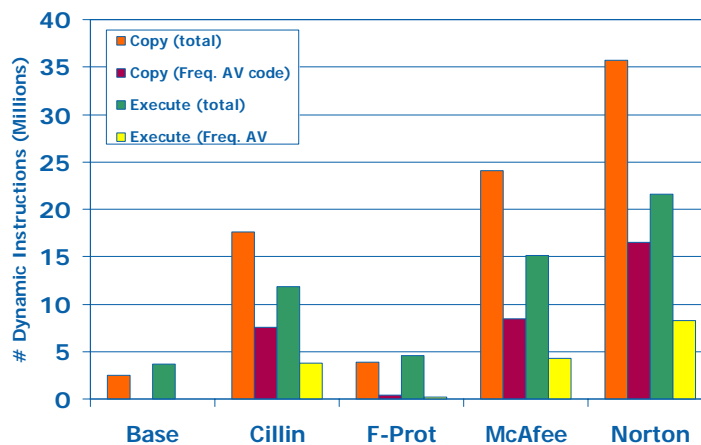
**1 hot code
trace makes up 17%
of the total execution**

```
mov esi,dword ptr 0x14[ebp]
test esi,esi
mov esi,dword ptr 0xf0[ebp]
mov edx,edi
movzx eax, dl
add eax,esi
movzx eax, byte ptr [eax]
cmp dword ptr 0xec[ebp],eax
ja 0x420cd84f
```

➤ A hot comparison block dominates execution

The Cost of Anti-virus Execution

Copy/Execution of "Hello, world" Application



Overhead Causes in Anti-Virus Security Mechanisms



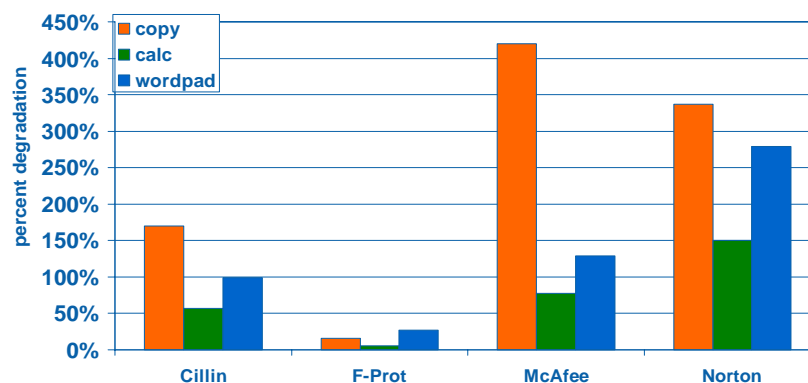
Signature Matching

- Program will refer to a dictionary of “signatures” or sequences of code known to be part of a malicious file
- If a signature is found in the file in question, it is marked as a virus
- Disadvantages:
 - Requires continuous updates
 - Cannot detect “zero-day attacks”

Heuristics

- Set of rules that the AV software will apply
 - For example, if the file contains self-modifying code
- If the file in question violates any of the given rules, it is marked as a virus
- Advantages:
 - May find virus variants
- Disadvantages:
 - Generates false positives

Antivirus Software Overhead



What's hot in AV applications?

Frequent "Hot" Code Examples



PC-Cillin

```
mov edx, dword ptr 0xb0[ebp]
inc ecx
add eax, 0xc
cmp ecx, edx
mov dword ptr 0xd4[ebp], ecx
jl 0xf45cc81a
```

F-Prot

```
mov ecx, dword ptr 0x8[ebp]
mov cl, byte ptr [ecx]
cmp cl, byte ptr 0xc[ebp]
je 0xf76a713
```

McAfee

```
xor edi, edi
mov ecx, dword ptr 0x8[ebp]
mov al, byte ptr 0x1[ebx]
lea edx, [edi][ecx]
mov cl, byte ptr [edi][ecx]
cmp al, cl
jne 0x1203c028
```

Norton


```
movzx edi, ax
imul edi, dword ptr 0xcfc[edx]
mov ebx, dword ptr 0x10[ebp]
add edi, ecx
cmp ebx, dword ptr [edi]
je 0xf6a13e02
```

➤ Frequent code exhibits similar structure

Instruction Footprint




Scenario: Copying a file from CDROM to the C drive



Workload Characterization

Copy file, 128 KB

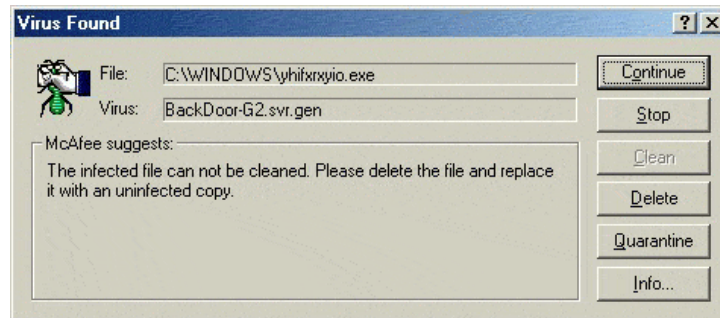


Architectural Extensions for SPAM Filtering and AV Scanning



- Can we exploit the characteristics of hot blocks and develop ISA extensions to accelerate these operations?
 - How will we address these issues when running in a virtualized environment? – Recent project with VMware
- ☺ Pin can help us identify answers to these challenging questions....

Pin for Information Flow Tracking



Zero-Day Attack Trends



In 2005, 74% of the top 50 code samples exposed confidential information

- Effects of Trojan Horses and Backdoors are more subtle

Symantec reported that 6 of the top 10 spyware were bundled with other programs

- Malware are executed without explicit consent

Zero-day attacks are increasing and are sold on black market

- Freshly authored malicious code can go undetected by even the most up-to-date virus scanners

We need a behavior-tracking mechanism that does not rely on known signatures

Example 1: PWSteal.Tarno.Q



Password-stealing Trojan Horse

1. Arrives in an email with a downloader in the attachment
Subject: Payment Receipt
Message: Dear Customer ...
Attachment: FILE.965658.exe
2. Downloads main part of the Trojan from a fixed location
[http://]dimmers.phpwebhosting.com/msupdate.exe?r=[UNIQUE_ID]
3. Creates a browser help object that runs every time Internet Explorer starts
4. Monitors windows and web pages with specific strings
gold, cash, bank, log, user, pin, memorable, secret
5. Periodically sends gather information using the following url:
[http://]neverdays.com/reporter.php

Example 2: Trojan.Lodeight.A



Trojan Horse that installs a Beagle and a backdoor

1. Contacts one of the following web sites using TCP port 80
[http://]www.tsaa.net/[REMOVED]
[http://]www.aureaorodeley.com/[REMOVED]
2. Downloads a remote file into the following folder and executes it. This remote file may be a mass-mailing worm such as W32.Beagle.CZ@mm.
%Windir%\[RANDOM NAME].exe
3. Opens a back door on TCP port 1084

Characteristics of Trojan Horses



	No User Intervention	Remotely Directed	Hard-coded Resources	Degrading Performance
PWSteal.Tarno.Q	✓		✓	
Trojan.Lodeight.A	✓	✓	✓	
Trojan.Vundo	✓		✓	✓
W32.MytoB.J@mm	✓	✓	✓	
Window-supdate.com	✓		✓	
W32/MyDoom.B	✓	✓	✓	
Phabot	✓	✓	✓	
Sendmail Trojan	✓		✓	
TCPWrappersTrojan	✓		✓	

105

Pin PLDI Tutorial 2007

Characteristics of Trojan Horses



- Malicious code is executed without user intervention
- Malicious code may be directed by a remote attacker once a connection is made
- Resources used by the malicious code (e.g. file names, URLs) are hard-coded in the binary
- Additional OS resources (processes, memory) are consumed by the malicious code
- How can we track this behavior dynamically?

➤Pin!!

106


Pin PLDI Tutorial 2007

Information Flow Tracking



Approach

- Track data sources and monitor information flow using Pin
- Send program behavior to back end whenever suspicious program behavior is suspected
- Provide analysis and policies to decide classify program behavior



Information Flow Tracking using Pin



- Pin tracks information flow in the program and identifies exact source of data
 - **USER_INPUT**: data is retrieved via user interaction
 - **FILE**: data is read from a file
 - **SOCKET**: data is retrieved from socket interface
 - **BINARY**: data is part of the program binary image
 - **HARDWARE**: data originated from hardware
- Pin maintains data source information for all memory locations and registers
- Propagates flow information by taking union of data sources of all operands

Example – Register Tracking



- We track flow from source to destination operands

...

%ebx - {}

%ecx - {BINARY1} /* ecx contains information from BINARY1 */

%edx - {SOCKET1} /* edx contains information from SOCKET1 */

%esi - {FILE2} /* esi contains information from FILE2 */

%edi - {}

...

- Assume the following XOR instruction:

xor %edx,%esi

which has the following semantics:

dst(%esi) := dst(%esi) XOR src(%edx)


- Pin will instrument this instruction and will insert an analysis routine to merge the source and destination operand information

%edx - {SOCKET1} /* edx contains information from SOCKET1 */
%esi - {SOCKET1, FILE2} /* esi contains information from FILE2 */

Information Flow Tracking using Pin



- **Different levels of abstraction**
- **Event Monitoring**
 - **Architectural Events**
 - Instructions executed
 - **OS Events**
 - System calls
 - **Library Events**
 - Library routines



Information Flow Tracking Prototype



System Calls

- Instrument selected system calls (12 in prototype)

Code Frequency

- Instrument every basic block
- Determine code "hotness"
- Application binary vs. shared object


Program Data Flow

- System call specific data flow
 - Tracking file loads, mapping memory to files ..
- Application data flow
 - Instrument memory access instructions
 - Instrument ALU instructions

111

Pin PLDI Tutorial 2007

Performance – Information Flow Tracking



112

Pin PLDI Tutorial 2007

Performance – Information Flow Tracking

Issues

- Significant overhead when considering both control and data flow (100x-1000x)
- Why is current implement underperforming?
 - Instrument every instruction for dataflow tracking
 - Could be done on a basic block level – register liveness
 - Track every basic block during hot/cold code analysis
 - Tracking path traces could reduce this overhead (Ball97)
- Pin probably not the best choice for full implementation due to runtime overhead – DynamoRIO (Duesterwald02)
- Consider adding hardware support – RIFLE (Vachharajani04)

☺ **Pin provides an effective mechanism for building a robust backend system and for exploring different information flow tracking schemes**

Using Pin in Security Research

• **Pin has been very useful in characterizing SPAM and Anti-virus workloads**

- ☺ Resulted in joint projects with VMWare and Network Engines

• **Pin has provided significant help in developing information flow tracking systems targeting zero-day attacks**

- ☺ Basis for a new startup company



Using Pin to Study Fault Tolerance and Program Behavior

Kim Hazelwood

David Kaeli

Dan Connors

Vijay Janapa Reddi

Pin-Based Fault Tolerance Analysis



Purpose:

- Simulate the occurrence of transient faults and analyze their impact on applications
- Construction of run-time system capable of providing software-centric fault tolerance service

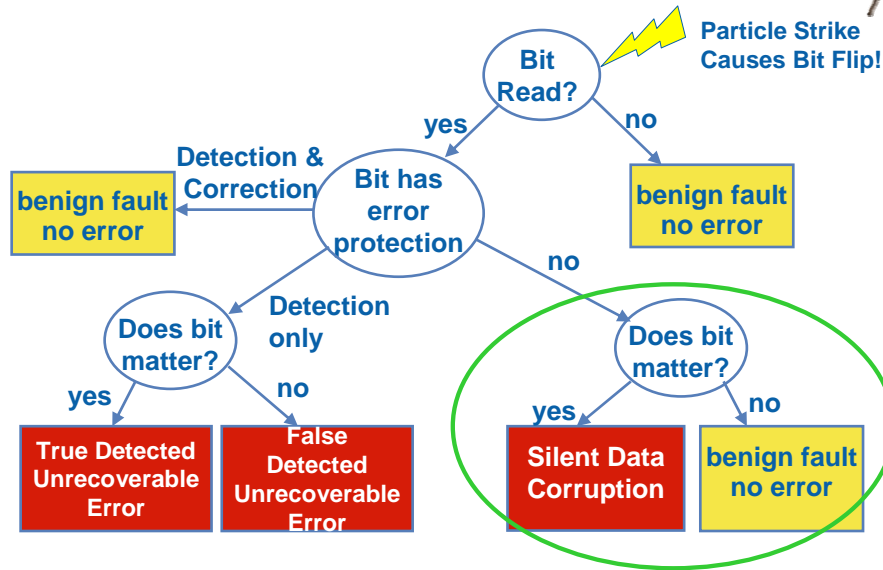
Pin

- Easy to model errors and the generation of faults and their impact
- Relatively fast (5-10 minutes per fault injection)
- Provides full program analysis

Research Work

- University of Colorado: Alex Shye, Joe Blomstedt, Harshad Sane, Alpesh Vaghasia, Tipp Moseley

Division of Transient Faults Analysis



117

Pin PLDI Tutorial 2007

Modeling Microarchitectural Faults in Pin


Accuracy of fault methodology depends on the complexity of the underlying system

- Microarchitecture, RTL, physical silicon

Build a microarchitectural model into Pin

- A low fidelity model may suffice
- Adds complexity and slows down simulation time

Emulate certain types of microarchitectural faults in Pin



118

Pin PLDI Tutorial 2007


Example: Destination/Source Register Transmission Fault



Fault occurs in latches when forwarding instruction output

Change architectural value of destination register at the instruction where fault occurs

NOTE: This is different than inserting fault into register file because the destination is selected based on the instruction where fault occurs



Example: Load Data Transmission Faults




Fault occurs when loading data from the memory system

Before load instruction, insert fault into memory

Execute load instruction

After load instruction, remove fault from memory (Cleanup)

NOTE: This models a fault occurring in the transmission of data from the STB or L1 Cache



Steps for Fault Analysis



Determine 'WHEN' the error occurs

Determine 'WHERE' the error occurs

Inject Error

Determine/Analyze Outcome

Step: WHEN



Sample Pin Tool: InstCount.C

- Purpose: Efficiently determines the number of dynamic instances of each static instruction

Output: For each static instruction

- Function name
- Dynamic instructions per static instruction

```
IP: 135000941 Count: 492714322 Func: propagate_block.104
IP: 135000939 Count: 492714322 Func: propagate_block.104
IP: 135000961 Count: 492701800 Func: propagate_block.104
IP: 135000959 Count: 492701800 Func: propagate_block.104
IP: 135000956 Count: 492701800 Func: propagate_block.104
IP: 135000950 Count: 492701800 Func: propagate_block.104
```

Step: WHEN



InstProf.C

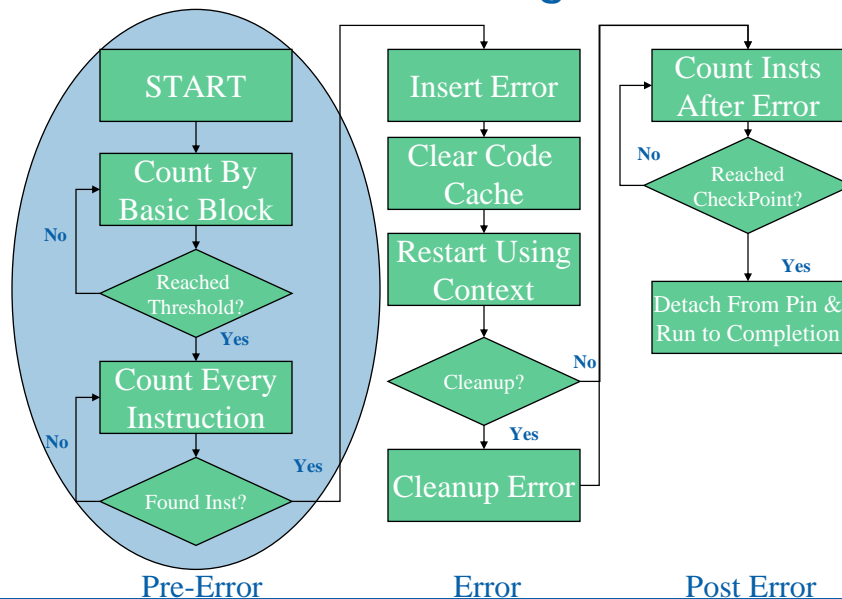
- Purpose: Traces basic blocks for contents and execution count

Output: For a program input

- Listing of dynamic block executions
- Used to generate a profile to select error injection point (opcode, function, etc)

```
BBL NumIns: 6    Count: 13356    Func: build_tree
804cb88 BINARY ADD    [Dest: ax] [Src: ax edx] MR: 1 MW: 0
804cb90 SHIFT  SHL    [Dest: eax] [Src: eax] MR: 0 MW: 0
804cb92 DATAXFER MOV  [Dest:]    [Src: esp edx ax] MR: 0 MW: 1
804cb97 BINARY INC    [Dest: edx] [Src: edx] MR: 0 MW: 0
804cb98 BINARY CMP    [Dest:]    [Src: edx] MR: 0 MW: 0
804cb9b COND_BR JLE   [Dest:]    [Src:] MR: 0 MW: 0
```

Error Insertion State Diagram



Step: WHERE

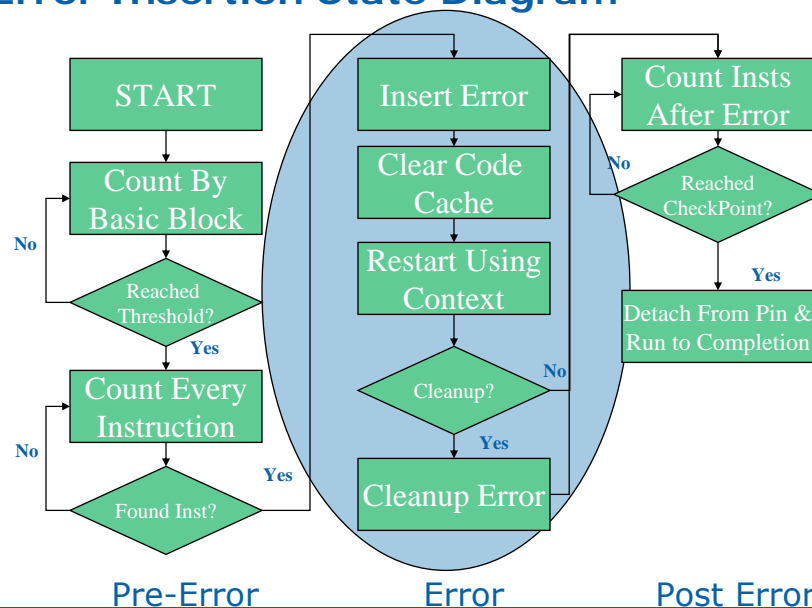
Reality:

- Where the transient fault occurs is a function of the size of the structure on the chip
- Faults can occur in both architectural and microarchitectural state

Approximation:

- Pin only provides architectural state, not microarchitectural state (no uops, for instance)
 - Either inject faults only into architectural state
 - Build an approximation for some microarchitectural state

Error Insertion State Diagram



Step: Injecting Error



Error Insertion Routine

```
VOID InsertFault(CONTEXT* _ctxt) {
    srand(curDynInst);
    GetFaultyBit(_ctxt, &faultReg, &faultBit);

    UINT32 old_val;  UINT32 new_val;
    old_val = PIN_GetContextReg(_ctxt, faultReg);
    faultMask = (1 << faultBit);
    new_val = old_val ^ faultMask;
    PIN_SetContextReg(_ctxt, faultReg, new_val);

    PIN_RemoveInstrumentation();
    faultDone = 1;
    PIN_ExecuteAt(_ctxt);
}
```

Step: Determining Outcome



Outcomes that can be tracked:

- Did the program complete?
- Did the program complete and have the correct IO result?
- If the program crashed, how many instructions were executed after fault injection before program crashed?
- If the program crashed, why did it crash (trapping signals)?

Register Fault Pin Tool: RegFault.C

MAIN

```

main(int argc, char * argv[]) {
    if (PIN_Init(argc, argv))
        return Usage();
    out_file.open(KnobOutputFile.Value().c_str());
    faultInst = KnobFaultInst.Value();

    TRACE_AddInstrumentFunction (Trace, 0);
    INS_AddInstrumentFunction(Instruction, 0);

    PIN_AddFiniFunction(Fini, 0);

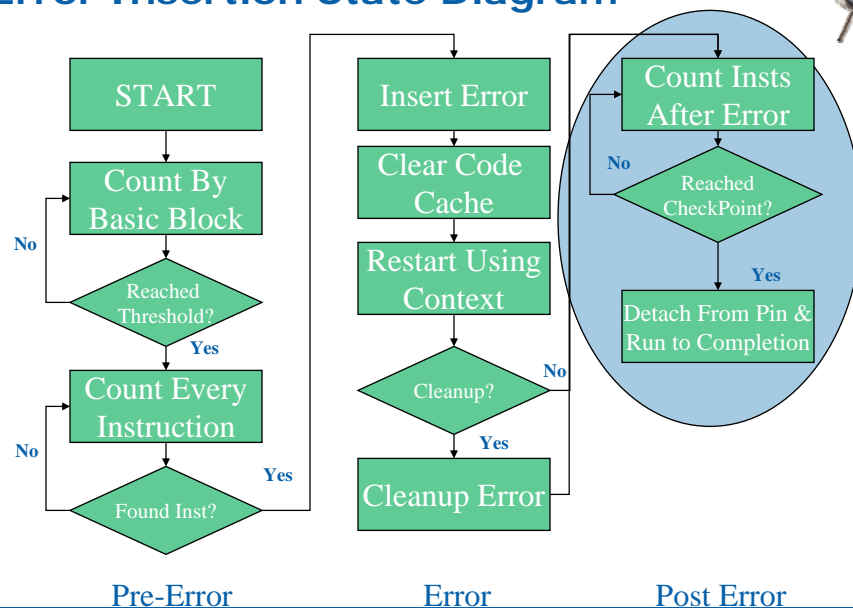
    PIN_AddSignalInterceptFunction(SIGSEGV, SigFunc, 0);
    PIN_AddSignalInterceptFunction(SIGFPE, SigFunc, 0);
    PIN_AddSignalInterceptFunction(SIGILL, SigFunc, 0);
    PIN_AddSignalInterceptFunction(SIGSYS, SigFunc, 0);

    PIN_StartProgram();
    return 0;
}
    
```

129

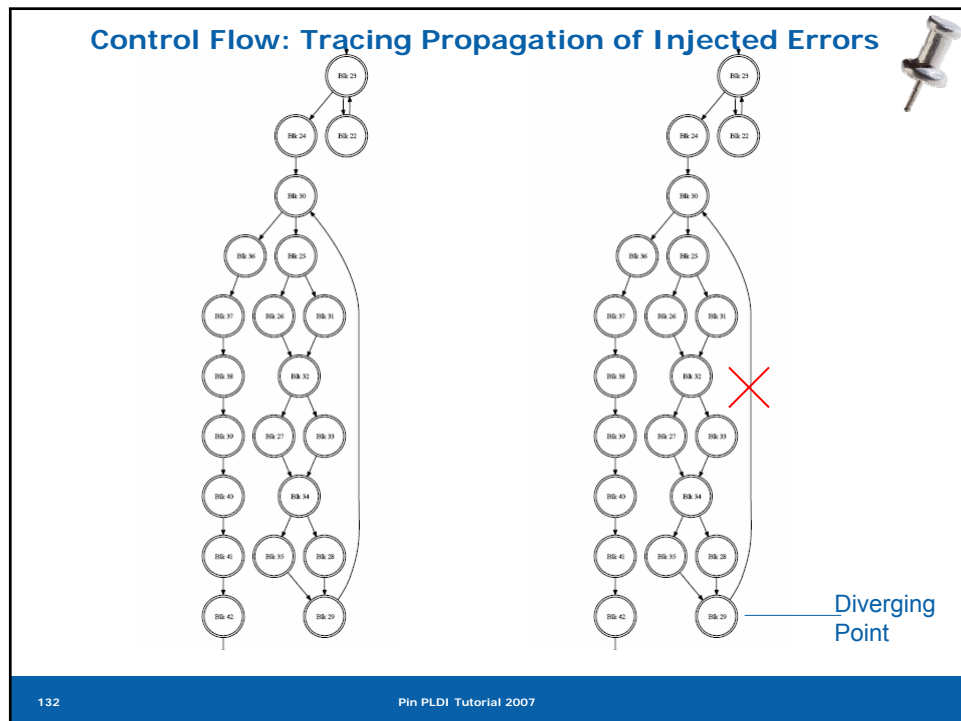
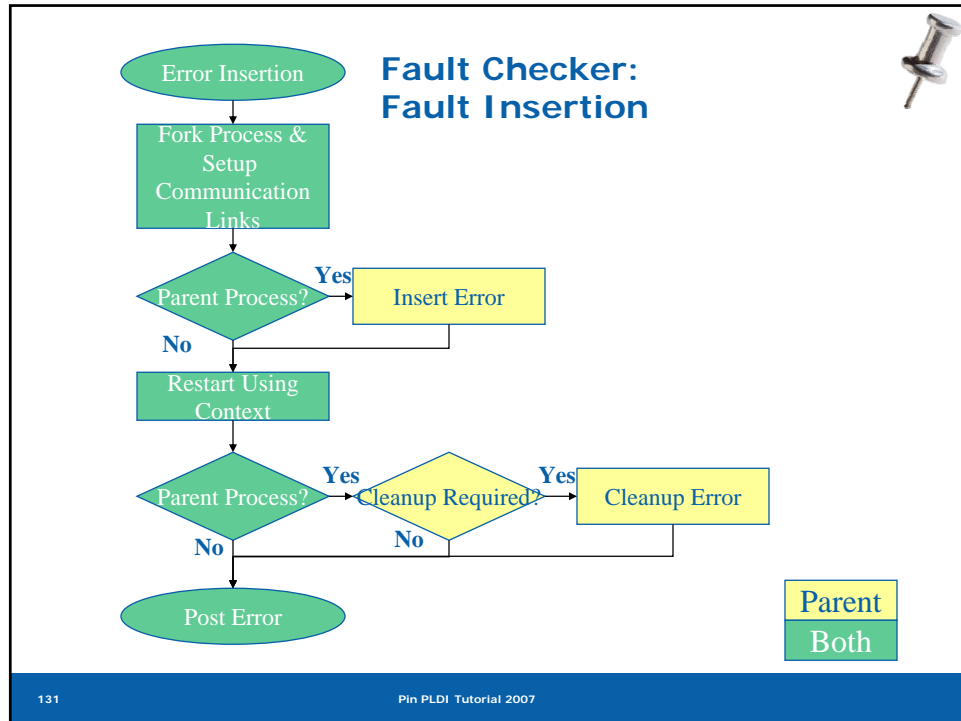
Pin PLDI Tutorial 2007

Error Insertion State Diagram




130

Pin PLDI Tutorial 2007




Data Flow: Tracing Propagation of Injected Errors



133

Pin PLDI Tutorial 2007

Fault Coverage Experimental Results



Watchdog timeout very rare so not shown


PLR detects all *Incorrect* and *Failed* cases

Effectively detects relevant faults and ignores benign faults

134

Pin PLDI Tutorial 2007


Function Analysis Experimental Results



164.gzi
Function

Per-function (top 10 function executed per application)

Fault Timeline Experimental Results



Error Injection until equal time segments of applications

Run-time System for Fault Tolerance



Process technology trends

- Single transistor error rate is expected to stay close to constant
- Number of transistors is increasing exponentially with each generation

Transient faults will be a problem for microprocessors!

Hardware Approaches


- Specialized redundant hardware, redundant multi-threading

Software Approaches


- Compiler solutions: instruction duplication, control flow checking
- Low-cost, flexible alternative but higher overhead

Goal: *Leverage available hardware parallelism in multi-core architectures to improve the performance of software-based transient fault tolerance*

Process-level Redundancy

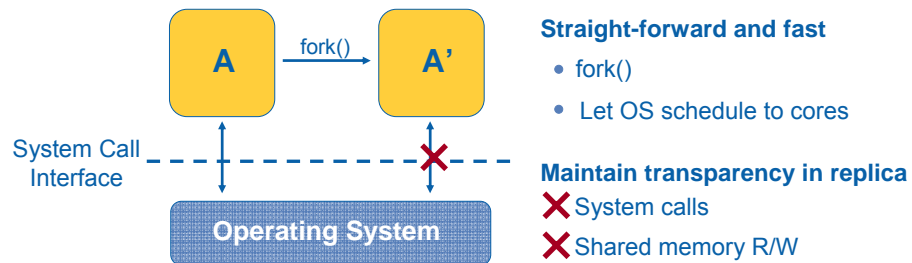


(a) Multiple Single Errors



(b) Multiple Simultaneous Errors

Replicating Processes



Replicas provide an extra copy of the program+input

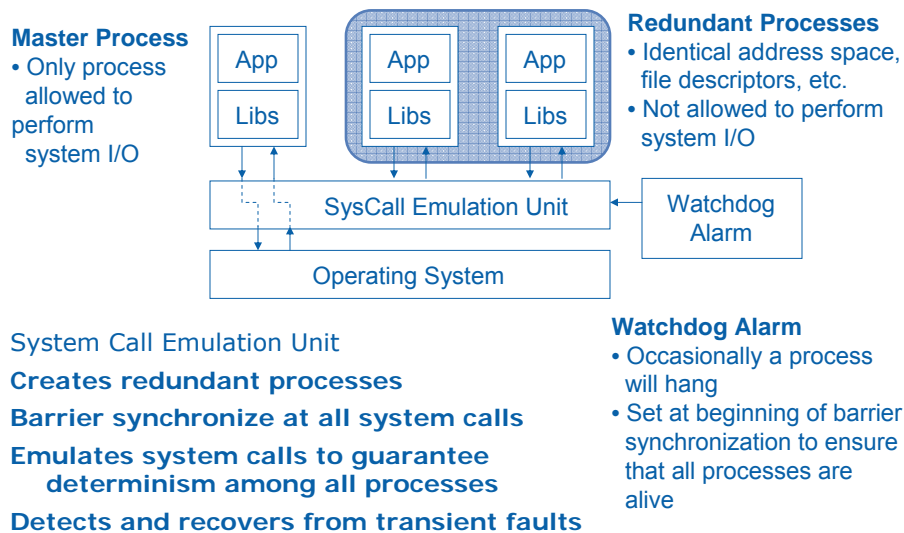
What can we do with this?

- Software transient fault tolerance
- Low-overhead program instrumentation
- More?

139

Pin PLDI Tutorial 2007


Process-Level Redundancy (PLR)



140

Pin PLDI Tutorial 2007

PLR Performance



Performance for single processor (PLR 1x1), 2 SMT processors (PLR 2x1) and 4 way SMP (PLR 4x1)

Slowdown for 4-way SMP only 1.26x

Conclusion



Fault insertion using Pin is a great way to determine the impacts faults have within an application

- Easy to use
- Enables full program analysis
- Accurately describes fault behavior once it has reached architectural state

Transient fault tolerance at 30% overhead

- Future work
 - Support non-determinism (shared memory, interrupts, multi-threading)
 - Fault coverage-performance trade-off in switching on/off



Pin-based Projects in Academia

Kim Hazelwood

David Kaeli

Dan Connors

Vijay Janapa Reddi

A Technique for Enabling & Supporting Field Failure Debugging



- **Problem**

In-house software quality is challenging, which results in field failures that are difficult to replicate and resolve

- **Approach**

Improve in-house debugging of field failures by

- (1) Recording & Replaying executions
- (2) Generating minimized executions for faster debugging

- **Who**

J. Clause and A. Orso @ Georgia Institute of Technology
ACM SIGSOFT Int'l. Conference on Software Engineering '07

Dytan: A Generic Dynamic Taint Analysis Framework



- **Problem**

Dynamic taint analysis is defined an adhoc-manner, which limits extendibility, experimentation & adaptability

- **Approach**

Define and develop a general framework that is customizable and performs data- and control-flow tainting

- **Who**

J. Clause, W. Li, A. Orso @ Georgia Institute of Technology
Int'l. Symposium on Software Testing and Analysis '07

Workload Characterization



- **Problem**

Extracting important trends from programs with large data sets is challenging

- **Approach**

Collect hardware-independent characteristics across program execution and apply them to statistical data analysis and machine learning techniques to find trends

- **Who**

K. Hoste and L. Eeckhout @ Ghent University

Loop-Centric Profiling



- **Problem**

Identifying parallelism is difficult

- **Approach**

Provide a hierarchical view of how much time is spent in loops, and the loops nested within them using (1) instrumentation and (2) light-weight sampling to automatically identify opportunities of parallelism

- **Who**

T. Moseley, D. Connors, D. Grunwald, R. Peri @
University of Colorado, Boulder and Intel Corporation
Int'l. Conference on Computing Frontiers (CF) '07

Shadow Profiling



- **Problem**

Attaining accurate profile information results in large overheads for runtime & feedback-directed optimizers

- **Approach**

`fork()` shadow copies of an application onto spare cores, which can be instrumented aggressively to collect accurate information without slowing the parent process

- **Who**

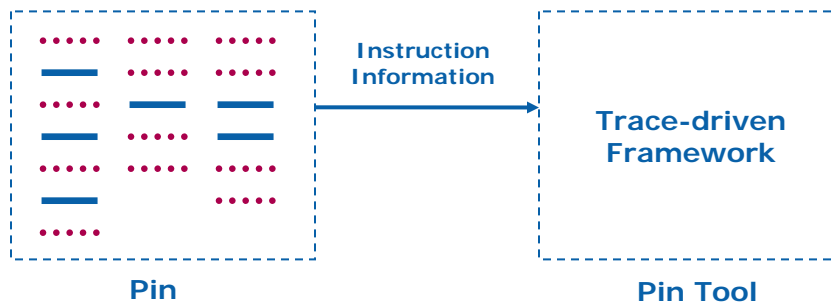
T. Moseley, A. Shye, V. J. Reddi, D. Grunwald, R. Peri
University of Colorado, Boulder and Intel Corporation
Int'l. Conference on Code Generation and Optimization (CGO) '07

Part Four: Exploratory Extensions and Hands-On Workshop

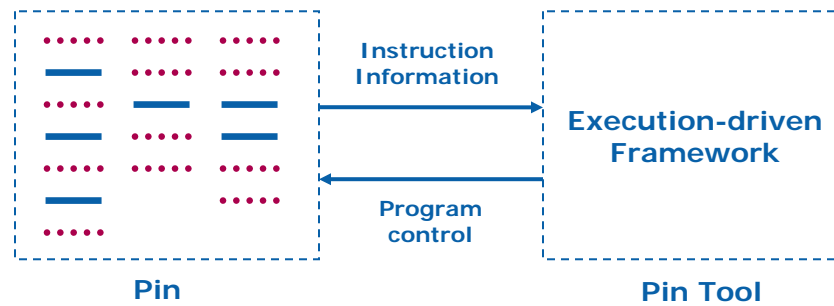
Kim Hazelwood
Dan Connors
David Kaeli

Vijay Janapa Reddi

Common use of Pin



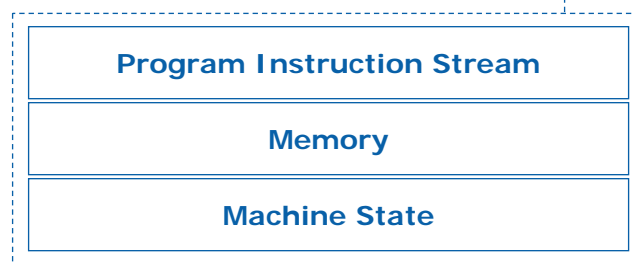
Driving execution using Pin



Session Objectives



- Building and Running Pin Tools
- Understanding **program execution** using Pin



- Putting it all together: Transactional Memory

Structure of a Pin Tool



```
FILE * trace;
```

Pin Tool traces Virtual Addresses

```
VOID RecordMemWrite(VOID * ip, VOID * va, UINT32 size) {  
    fprintf(trace,"%p: W %p %d\n", ip, va, size);  
}  
  
VOID Instruction(INS ins, VOID *v) {  
    if (INS_IsMemoryWrite(ins)) {  
        INS_InsertCall(ins, IPOINT_BEFORE, AFUNPTR(RecordMemWrite),  
                        IARG_INST_PTR,  
                        IARG_MEMORYWRITE_VA,  
                        IARG_MEMORYWRITE_SIZE, IARG_END);  
    }  
}
```

Analysis

```
int main(int argc, char *argv[]) {  
    PIN_Init(argc, argv);  
    trace = fopen("atrace.out", "w");  
    INS_AddInstrumentFunction(Instruction, 0);  
    PIN_StartProgram();  
    return 0;  
}
```

Instrumentation

Callback Registration



Architectural State Interposition



- **Observe instruction operands and their values**
 - IARG_BRANCH_TAKEN, IARG_REG_VALUE, IARG_CONTEXT, ...
- **Modify register values**
- **Save and restore state**
- **Instruction emulation**




Modify architectural state



- **Alter register values via instrumentation**

- IARG_REG_REFERENCE <register>
- PIN_REGISTER *



```
/* ===== Instrumentation routine ===== */
if (INS_IsRDTSC(ins))
{
    INS_InsertCall(ins, IPOINT_AFTER,
        (AFUNPTR) DeterministicRDTSC,
        IARG_REG_REFERENCE, REG_EDX,
        IARG_REG_REFERENCE, REG_EAX,
        IARG_END);
}

/* ===== Analysis routine ===== */
VOID DeterministicRDTSC(ADDRINT *pEDX, ADDRINT *pEAX)
{
    static UINT64 _edx_eax = 0;
    _edx_eax += 1;

    *pEDX = (_edx_eax & 0xffffffff00000000ULL) >> 32;
    *pEAX = _edx_eax & 0x00000000ffffffffULL;
}
```

155

Pin PLDI Tutorial 2007




Save and Resume Execution



- **Capture snapshots of the machine state to resume at a later point**

- IARG_CHECKPOINT
- PIN_SaveCheckpoint(CHECKPOINT *, CHECKPOINT *)
- PIN_Resume(CHECKPOINT *)



156

Pin PLDI Tutorial 2007



Save and Resume Execution (2)



- **IARG_CHECKPOINT**
 - Pin generates a snapshot (includes instrumented state)
- **PIN_SaveCheckpoint** (CHECKPOINT *src, CHECKPOINT *dst)
 - Extract and copy state from handle(src) to local buffer(dst)



Pin
Stream

```
/* ===== Instrumentation routine ===== */
INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) Save,
               IARG_CHECKPOINT,
               IARG_END);
```

```
/* ===== Analysis routine ===== */
CHECKPOINT ckpt;
```


```
VOID Save(CHECKPOINT* _ckpt)
{
    PIN_SaveCheckpoint(_ckpt, &ckpt);
}
```



Save and Resume Execution (3)



- **PIN_Resume(CHECKPOINT *)**
 - Restore processor state to saved checkpoint
 - Continue execution



Pin
Stream

```
/* ===== Instrumentation routine ===== */
INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) Back,
               IARG_END);
```

```
/* ===== Analysis routine ===== */
CHECKPOINT ckpt;
```

```
VOID Back()
{
    PIN_Resume(&ckpt);

    assert(false); /* PIN_Resume does not return! */
}
```

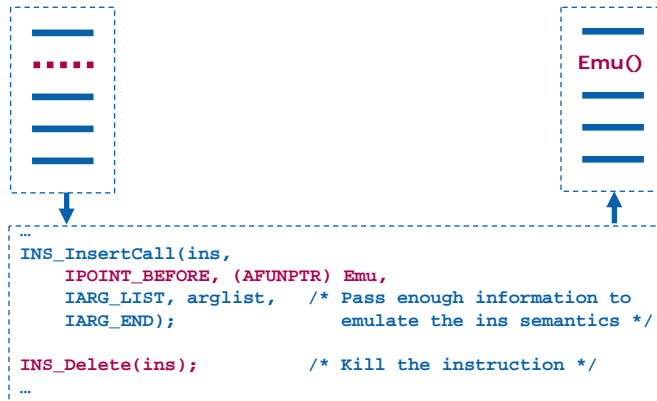



Instruction Emulation



- Emulate the semantics of (new) instructions

- (1) Locate emu instruction
- (2) Marshall semantics
- (3) Substitute emu function
- (4) Delete emu instruction



Emulating a Load Instruction



```
#include "pin.H"
#include "pin_isa.H"

ADDRINT DoLoad(REG reg, ADDRINT * addr) {
    return *addr;
}

VOID EmulateLoad(INS ins, VOID* v) {
    if (INS_Opcode(ins) == XEDICLASS_MOV && INS_IsMemoryRead(ins) &&
        INS_OperandIsReg(ins, 0) && INS_OperandIsMemory(ins, 1)) {
        INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) DoLoad,
            IARG_UINT32, REG(INS_OperandReg(ins, 0)),
            IARG_MEMORYREAD_EA,
            IARG_RETURN_REGS, INS_OperandReg(ins, 0),
            IARG_END);

        INS_Delete(ins);
    }
}

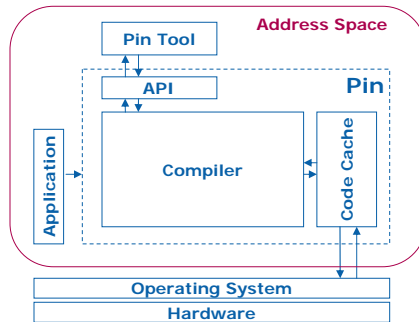
/* Emulate load type */
op0 <- *op1

void main(int argc, char * argv[]) {
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(EmulateLoad, 0);
    PIN_StartProgram();
}
```

Memory Behavior



- **Memory access tracing**
 - IARG_MEMORYREAD_EA, IARG_MEMORYWRITE_EA, ...
- **Modify program memory**
 - Pin Tool resides in the process' address space

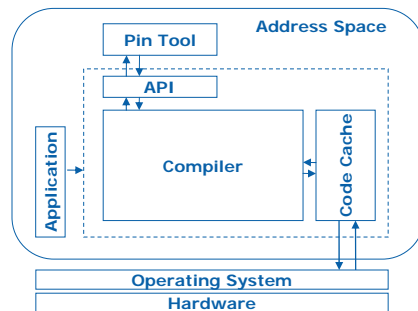


Change memory directly
(*addr = 0x123)

Controlling Program Execution

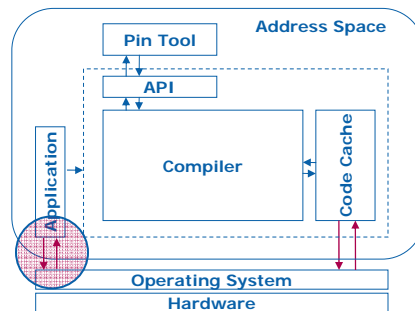


Pin (JIT)



Only translated code cached in the *Code Cache* is executed
Pros : Complete coverage
Cons: Slow

Pin (Probes)




Original code, and translated code are executed intermixed with one another
Pros : Fast
Cons: Limited coverage

Program

Executing @ Arbitrary Locations



- **JIT-mode** (execute only translated code)
 - IARG_CONTEXT
 - PIN_ExecuteAt (CONTEXT *)



163


Pin PLDI Tutorial 2007

Program

Executing @ Arbitrary Locations (2)



- **IARG_CONTEXT**
 - Pin generates program's perception of machine state
- **PIN_ExecuteAt** (CONTEXT *)
 - Continue executing at context state



```

/* ===== Instrumentation routine ===== */
if (INS_Address(ins) == 0x40000000 /* Foo: */)
  INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR) Jump2Bar,
                IARG_CONTEXT,
                IARG_END);

/* ===== Analysis routine ===== */
VOID Jump2Bar(CONTEXT *ctxt)
{
  PIN_SetContextReg(ctxt, REG_INST_PTR, Bar);
  PIN_ExecuteAt(ctxt);
  assert(false); /* PIN_ExecuteAt does not return! */
}

```

164

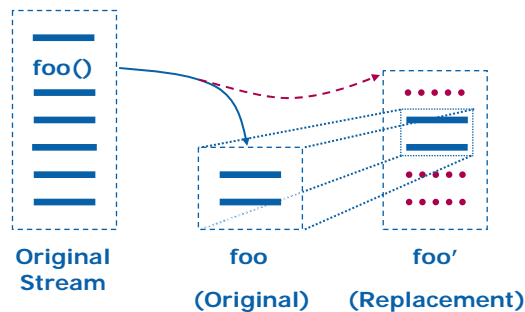
Pin PLDI Tutorial 2007

Program

Changing Program Code (Probe-mode)



- **PIN_ReplaceProbed** (RTN, AFUNPTR)
 - Redirect control flow to new functions in the Pin Tool
- **PIN_ReplaceSignatureProbed** (RTN, AFUNPTR, ...)
 - (1) Redirect control flow (2) Rewrite function prototypes
 - (3) Use Pin arguments (IARG's)



165

Pin PLDI Tutorial 2007

Program

Replacing malloc() in Application



```
typedef VOID * (*FUNCPTR_MALLOC)(size_t);

VOID * MyMalloc(FUNCPTR_MALLOC orgMalloc, UINT32 size, ADDRINT returnIp) {
    FUNCPTR_MALLOC poolMalloc = LookupMallocPool(returnIp, size);

    return (poolMalloc) ? poolMalloc(size) : orgMalloc(size);
}

VOID ImageLoad(IMG img, VOID *v) {
    RTN mallocRTN = RTN_FindByName(img, "malloc");

    if (RTN_Valid(rtn)) {
        PROTO prototype = PROTO_Allocate(PIN_PARG(void *), CALLINGSTD_CDECL,
                                          "malloc", PIN_PARG(int), PIN_PARG_END());

        RTN_ReplaceSignatureProbed(mallocRTN, (AFUNPTR) MyMalloc,
                                   IARG_PROTOTYPE, prototype,           /* Function prototype */
                                   IARG_ORIG_FUNCPTR,                   /* Handle to application's malloc */
                                   IARG_FUNCARG_ENTRYPOINT_VALUE, 0,     /* First argument to malloc */
                                   IARG_RETURN_IP,                  /* IP of caller */
                                   IARG_END);

        PROTO_Free( proto_malloc );
    }
}
```

166

Pin PLDI Tutorial 2007



Source-level Probing



- Instrument only specific regions of the source

```
#include <stdio.h>
#include "pinapp.h"

int a[10];
int main()
{
    void * th = PIN_NewThread();

    printf("Thread handle %p\n", th);


    PIN_ExecuteInstrumented(th);

    for (int i = 0; i < 10; i++)
    {
        a[i] = i;
    }


    PIN_ExecuteUninstrumented();
    return 0;
}
```

Pin

Pin
Tool




Putting it all together: TMM




Transactional Memory Model


- Checkpoint architectural and memory state
- Log memory values modified by transaction
- Verify conflicts across parallel transactions
- Commit or Abort active transaction




Transactional Memory Model (1)




```
/* === Instrumentation routine === */  
  
if (RTN_Address(rtn) == XBEGIN)  
{  
    RTN_InsertCall(rtn, IPOINT_BEFORE,  
        AFUNPTR(BeginTransaction),  
        IARG_THREAD_ID,  
        IARG_CHECKPOINT,  
        IARG_END);  
}  
  
/* ===== Analysis routine ===== */  
  
CHECKPOINT chkpt[NTHREADS];  
  
void BeginTransaction(int tid,  
    CHECKPOINT *_chkpt)  
{  
    PIN_SaveCheckpoint(_chkpt,  
        chkpt[tid];  
}
```




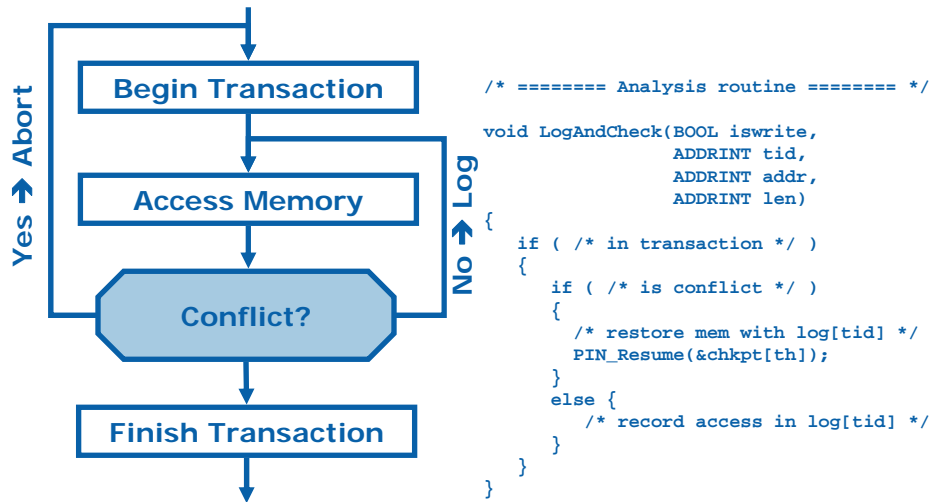
Transactional Memory Model (2)



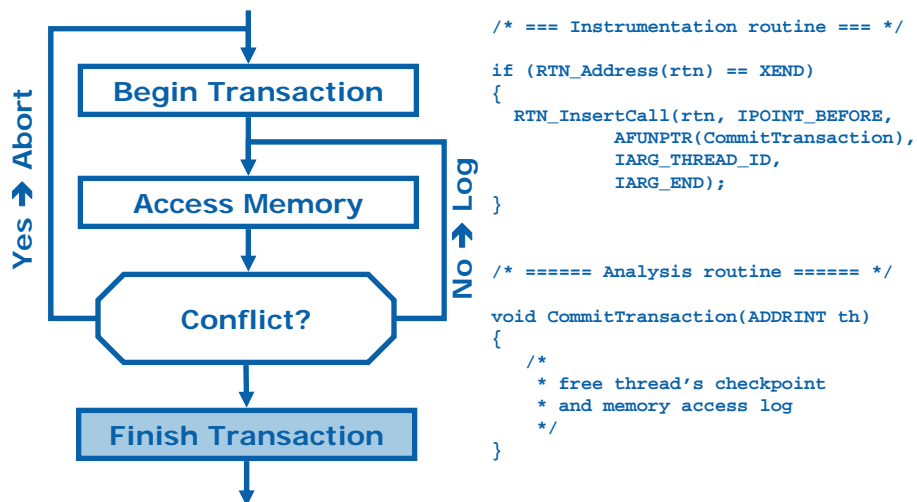
```
/* ===== Instrumentation routine ===== */  
  
void Instruction(INS ins, void *v)  
{  
    if (INS_IsMemoryWrite(ins))  
        INS_InsertCall(ins, IPOINT_BEFORE,  
            (AFUNPTR) LogAndCheck,  
            IARG_BOOL, true,  
            IARG_THREAD_ID,  
            IARG_MEMORYWRITE_EA,  
            IARG_MEMORYWRITE_SIZE,  
            IARG_END);  
  
    if (INS_IsMemoryRead(ins))  
        INS_InsertCall(ins, IPOINT_BEFORE,  
            (AFUNPTR) LogAndCheck,  
            IARG_BOOL, false,  
            IARG_THREAD_ID,  
            IARG_MEMORYREAD_EA,  
            IARG_MEMORYREAD_SIZE,  
            IARG_END);  
}
```




Transactional Memory Model (3)



Transactional Memory Model (4)






Demo of Transactional Memory



Multi-threaded Application

Transactional Memory Pin Tool



T1

```
XBEGIN();
for (uint32_t i = 0; i < MAX; i++)
{
    myarray[i] = 1;
}
XEND();
```

T2

```
XBEGIN();
for (int32_t i = MAX-1; i >= 0; i--)
{
    myarray[i] = 2;
}
XEND();
```

Yes → Abort

Begin Transaction


Access Memory

Conflict?


Finish Transaction

No → Log

Pin (user-level) → PinOS (system-level)



Pin



Pin → Pin the OS!

PinOS: A Programmable Framework for Whole-System Dynamic Instrumentation.
Prashanth P. Bungale, C. K. Luk. *Proceedings of Virtual Execution Environments (VEE 2007)*

Trace **Physical** and Virtual Addresses



```
FILE * trace;

VOID RecordMemWrite(VOID * ip, VOID * va, VOID * pa, UINT32 size) {
    Host_fprintf(trace,"%p: W %p %p %d\n", ip, va, pa, size);
}

VOID Instruction(INS ins, VOID *v) {
    if (INS_IsMemoryWrite(ins)) {
        INS_InsertCall(ins, IPOINT_BEFORE, AFUNPTR(RecordMemWrite),
            IARG_INST_PTR,
            IARG_MEMORYWRITE_VA,
            IARG_MEMORYWRITE_PA,
            IARG_MEMORYWRITE_SIZE, IARG_END);
    }
}

int main(int argc, char *argv[]) {
    PIN_Init(argc, argv);
    trace = Host_fopen("atrace.out", "w");
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_StartProgram();
    return 0;
}
```

**PinOS requires
minimal API changes**

175

Pin PLDI Tutorial 2007

Concluding Remarks



- **Dynamic instrumentation framework (Free!)**
 - Transparent across platforms and environments
 - Platforms: IA32, EM64T, Itanium, and Xscale
 - Operating Systems: Linux, Windows, MacOS
- **Sample tools** (use as templates)
 - Cache simulators, Branch predictors, Memory checkers, Instruction and Memory tracing, Profiling, Sampling ...
- **Write your own tools!**
- **Visit us @ <http://rogue.colorado.edu/wikipin>**

176

Pin PLDI Tutorial 2007

NOTES

NOTES